

Example: $B := LB$ – Team: Robert van de Geijn

1.1 Operation

Consider the operation

$$B := LB$$

where L is a $m \times m$ lower triangular matrix and B is a $m \times n$ matrix. This is a special case of triangular matrix-matrix multiplication, with the Lower triangular matrix on the Left, and the triangular matrix is Not transposed. We will refer to this operation as TRMM_LLNN where the LLNN stands for left lower no-transpose nonunit diagonal. The Nonunit diagonal means we will use the entries of the matrix that are stored on the diagonal. In some other cases, the entries on the diagonal may implicitly be taken to all equal one (Unit).

1.2 Precondition and postcondition

In the precondition

$$B = \hat{B}$$

\hat{B} denotes the original contents of B . This allows us to express the state upon completion, the postcondition, as

$$B = L\hat{B}.$$

It is implicitly assumed that L is nonunit lower triangular.

1.3 Partitioned Matrix Expressions and loop invariants

There are two PME's for this operation.

1.3.1 PME 1

To derive the second PME, partition

$$L \rightarrow \left(\begin{array}{c|c} L_{TL} & 0 \\ \hline L_{BL} & L_{BR} \end{array} \right), \quad \text{and} \quad B \rightarrow \left(\begin{array}{c} B_T \\ \hline B_B \end{array} \right).$$

Substituting these into the postcondition yields

$$\begin{pmatrix} B_T \\ B_B \end{pmatrix} = \begin{pmatrix} L_{TL} & 0 \\ L_{BL} & L_{BR} \end{pmatrix} \begin{pmatrix} \hat{B}_T \\ \hat{B}_B \end{pmatrix}$$

or, equivalently,

$$\begin{pmatrix} B_T \\ B_B \end{pmatrix} = \begin{pmatrix} L_{TL}\hat{B}_T \\ L_{BL}\hat{B}_T + L_{BR}\hat{B}_B \end{pmatrix},$$

which we refer to as the first PME for this operations.

From this, we can choose two loop invariants:

Invariant 1: $\begin{pmatrix} B_T \\ B_B \end{pmatrix} = \begin{pmatrix} \hat{B}_T \\ L_{BR}\hat{B}_B \end{pmatrix}.$

(The top part has been left alone and the bottom part has been partially computed).

Invariant 2: $\begin{pmatrix} B_T \\ B_B \end{pmatrix} = \begin{pmatrix} \hat{B}_T \\ L_{BL}\hat{B}_T + L_{BR}\hat{B}_B \end{pmatrix}.$

(The top part has been left alone and the bottom part has been completely computed).

1.3.2 PME 2

To derive the second PME, partition

$$B \rightarrow \begin{pmatrix} B_L & B_R \end{pmatrix}$$

and do not partition L . Substituting these into the postcondition yields

$$\begin{pmatrix} B_L & B_R \end{pmatrix} = L \begin{pmatrix} \hat{B}_L & \hat{B}_R \end{pmatrix}$$

or, equivalently,

$$\begin{pmatrix} B_L & B_R \end{pmatrix} = \begin{pmatrix} L\hat{B}_L & L\hat{B}_R \end{pmatrix},$$

which we refer to as the second PME.

From this, we can choose two more loop invariants:

Invariant 3: $\begin{pmatrix} B_L & B_R \end{pmatrix} = \begin{pmatrix} L\hat{B}_L & \hat{B}_R \end{pmatrix}.$

(The left part has been completely finished and the right part has been left untouched).

Invariant 4: $\begin{pmatrix} B_L & B_R \end{pmatrix} = \begin{pmatrix} \hat{B}_L & L\hat{B}_R \end{pmatrix}.$

(The left part has been completely finished and the right part has been left untouched).

1.3.3 Notes

How do I decide to partition the matrices in the postcondition?

- Pick a matrix (operand), any matrix.
- If that matrix has
 - a triangular structure (in storage), then you want to either partition it into four quadrants, or not at all. Symmetric matrices and triangular matrices have a triangular structure (in storage).
 - no particular structure, then you partition it vertically (left-right), horizontally (top-bottom), or not at all.
- Next, partition the other matrices similarly, but conformally (meaning the resulting multiplications with the parts are legal).

Take our problem here: $B := LB$. Start by partitioning L into quadrants:

$$B = \left(\begin{array}{c|c} L_{TL} & 0 \\ \hline L_{BL} & L_{BR} \end{array} \right) \widehat{B}.$$

Now, the way partitioned matrix multiplication works, this doesn't make sense:

$$B = \underbrace{\left(\begin{array}{c|c} L_{TL} & 0 \\ \hline L_{BL} & L_{BR} \end{array} \right) \widehat{B}}_{\left(\begin{array}{c} L_{TL} \times \text{something} + 0 \times \text{something} \\ \hline L_{BL} \times \text{something} + L_{BR} \times \text{something} \end{array} \right)}.$$

So, we need to also partition B into a top part and a bottom part:

$$\left(\begin{array}{c} B_T \\ \hline B_B \end{array} \right) = \underbrace{\left(\begin{array}{c|c} L_{TL} & 0 \\ \hline L_{BL} & L_{BR} \end{array} \right) \left(\begin{array}{c} \widehat{B}_T \\ \hline \widehat{B}_B \end{array} \right)}_{\left(\begin{array}{c} L_{TL}\widehat{B}_T + 0 \times \widehat{B}_B \\ \hline L_{BL}\widehat{B}_T + L_{BR}\widehat{B}_B \end{array} \right)}.$$

Alternatively, what if you don't partition L ? You have to partition *something* so let's try partitioning B :

$$\left(\begin{array}{c} B_T \\ \hline B_B \end{array} \right) = L \left(\begin{array}{c} \widehat{B}_T \\ \hline \widehat{B}_B \end{array} \right).$$


But that doesn't work... Instead

$$\left(B_L \mid B_R \right) = L \left(\widehat{B}_L \mid \widehat{B}_R \right) = \left(L\widehat{B}_L \mid L\widehat{B}_R \right)$$

works just fine.

1.4 Deriving all unblocked algorithms

The below table summarizes all loop invariants, with links to all files related to this operation.

The worksheet and code skeletons were generated using the  [Spark webpage](#).

	Invariant	Derivations	Implementations
1	$\left(\frac{B_T}{B_B} \right) = \left(\frac{\hat{B}_T}{L_{BR}\hat{B}_B} \right)$	PDF	trmm_llnn_unb_var1.mlx trmm_llnn_unb_var1.c
2	$\left(\frac{B_T}{B_B} \right) = \left(\frac{\hat{B}_T}{L_{BL}B_T + L_{BR}\hat{B}_B} \right)$	PDF	trmm_llnn_unb_var2.mlx trmm_llnn_unb_var2.c
3	$\left(B_L \mid B_R \right) = \left(L\hat{B}_L \mid \hat{B}_R \right)$	PDF	trmm_llnn_unb_var3.mlx trmm_llnn_unb_var3.c
4	$\left(B_L \mid B_R \right) = \left(\hat{B}_L \mid L\hat{B}_R \right)$	PDF	trmm_llnn_unb_var4.mlx trmm_llnn_unb_var4.c